

# ASSIGNMENT

Name: - Aryan Omkar Ashar

Registration Number: - 19BAI10094

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline

import tensorflow as tf
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.utils import to_categorical

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Conv2D, MaxPool2D, Flatten, Dropout, BatchNormalization
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.preprocessing.image import ImageDataGenerator

from sklearn.metrics import ConfusionMatrixDisplay
from sklearn.metrics import classification_report, confusion_matrix
```

## DOWNLOADING THE DATASET

In [2]:

```
(X_train, y_train), (X_test, y_test) = cifar10.load_data()
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>  
170500096/170498071 [=====] - 3s 0us/step

In [3]:

```
print(f"X_train shape: {X_train.shape}")
print(f"y_train shape: {y_train.shape}")
print(f"X_test shape: {X_test.shape}")
print(f"y_test shape: {y_test.shape}")
```


```
X_train shape: (50000, 32, 32, 3)
y_train shape: (50000, 1)
X_test shape: (10000, 32, 32, 3)
y_test shape: (10000, 1)
```

## Data Visualization

In [4]:

```
labels = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
W_grid = 10
L_grid = 10
fig, axes = plt.subplots(L_grid, W_grid, figsize = (17,17))
axes = axes.ravel()
n_train = len(X_train)
for i in np.arange(0, W_grid * L_grid): # create evenly spaces variables
    index = np.random.randint(0, n_train)
    axes[i].imshow(X_train[index,1:])
    label_index = int(y_train[index])
    axes[i].set_title(labels[label_index], fontsize = 8)
    axes[i].axis('off')
```

```
plt.subplots_adjust(hspace=0.4)
```




In [5]:


```
classes_name = ['Airplane', 'Automobile', 'Bird', 'Cat', 'Deer', 'Dog', 'Frog', 'Horse',
                'Ship', 'Truck']

classes, counts = np.unique(y_train, return_counts=True)
plt.barh(classes_name, counts)
plt.title('Class distribution in training set')
```

Out[5]:

Text(0.5, 1.0, 'Class distribution in training set')






In [6]:

```
classes, counts = np.unique(y_test, return_counts=True)
plt.barh(classes_name, counts)
plt.title('Class distribution in testing set')
```

Out[6]:

Text(0.5, 1.0, 'Class distribution in testing set')



**The class are equally distributed**

## Data Preprocessing

In [7]:

```
# Scale the data
X_train = X_train / 255.0
X_test = X_test / 255.0

# Transform target variable into one-hotencoding
y_cat_train = to_categorical(y_train, 10)
y_cat_test = to_categorical(y_test, 10)
```

In [8]:

y\_cat\_train

Out[8]:

```
array([[0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 1.],
       [0., 0., 0., ..., 0., 0., 1.],
       ...,
       [0., 0., 0., ..., 0., 0., 1.],
       [0., 1., 0., ..., 0., 0., 0.],
       [0., 1., 0., ..., 0., 0., 0.]], dtype=float32)
```

## BUILDING THE MODEL

In [9]:

```
model = Sequential()

# Convolutional Layer
model.add(Conv2D(filters=32, kernel_size=(3, 3), input_shape=(32, 32, 3), activation='re
```

```

lu', padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(filters=32, kernel_size=(3, 3), input_shape=(32, 32, 3), activation='relu', padding='same'))
model.add(BatchNormalization())
# Pooling layer
model.add(MaxPool2D(pool_size=(2, 2)))
# Dropout layers
model.add(Dropout(0.25))

model.add(Conv2D(filters=64, kernel_size=(3, 3), input_shape=(32, 32, 3), activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(filters=64, kernel_size=(3, 3), input_shape=(32, 32, 3), activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(MaxPool2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(filters=128, kernel_size=(3, 3), input_shape=(32, 32, 3), activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(filters=128, kernel_size=(3, 3), input_shape=(32, 32, 3), activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(MaxPool2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())
# model.add(Dropout(0.2))
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.25))
model.add(Dense(10, activation='softmax'))

METRICS = [
    'accuracy',
    tf.keras.metrics.Precision(name='precision'),
    tf.keras.metrics.Recall(name='recall')
]
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=METRICS)

```

In [10]:

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
batch_normalization (BatchNo	(None, 32, 32, 32)	128
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
batch_normalization_1 (Batch	(None, 32, 32, 32)	128
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
dropout (Dropout)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18496
batch_normalization_2 (Batch	(None, 16, 16, 64)	256
conv2d_3 (Conv2D)	(None, 16, 16, 64)	36928
batch_normalization_3 (Batch	(None, 16, 16, 64)	256
max_pooling2d_1 (MaxPooling2	(None, 8, 8, 64)	0
dropout_1 (Dropout)	(None, 8, 8, 64)	0

conv2d_4 (Conv2D)	(None, 8, 8, 128)	73856
batch_normalization_4 (Batch Normalization)	(None, 8, 8, 128)	512
conv2d_5 (Conv2D)	(None, 8, 8, 128)	147584
batch_normalization_5 (Batch Normalization)	(None, 8, 8, 128)	512
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 128)	0
dropout_2 (Dropout)	(None, 4, 4, 128)	0
flatten (Flatten)	(None, 2048)	0
dense (Dense)	(None, 128)	262272
dropout_3 (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 10)	1290
=====		
Total params: 552,362		
Trainable params: 551,466		
Non-trainable params: 896		

## Early Stopping

In [11]:

```
early_stop = EarlyStopping(monitor='val_loss', patience=2)
```

## Data Augmentations and Model Training

In [12]:

```
batch_size = 32
data_generator = ImageDataGenerator(width_shift_range=0.1, height_shift_range=0.1, horizontal_flip=True)
train_generator = data_generator.flow(X_train, y_cat_train, batch_size)
steps_per_epoch = X_train.shape[0] // batch_size
```

In [13]:

```
r = model.fit(train_generator,
              epochs=50,
              steps_per_epoch=steps_per_epoch,
              validation_data=(X_test, y_cat_test),
              #         callbacks=[early_stop],
              #         batch_size=batch_size,
              )
```

Epoch 1/50

1562/1562 [=====] - 40s 20ms/step - loss: 1.9058 - accuracy: 0.3316 - precision: 0.5119 - recall: 0.1173 - val\_loss: 1.3044 - val\_accuracy: 0.5370 - val\_precision: 0.7033 - val\_recall: 0.3861

Epoch 2/50

1562/1562 [=====] - 31s 20ms/step - loss: 1.2966 - accuracy: 0.5407 - precision: 0.7096 - recall: 0.3617 - val\_loss: 0.9729 - val\_accuracy: 0.6640 - val\_precision: 0.7745 - val\_recall: 0.5560

Epoch 3/50

1562/1562 [=====] - 31s 20ms/step - loss: 1.0885 - accuracy: 0.6217 - precision: 0.7622 - recall: 0.4823 - val\_loss: 0.9059 - val\_accuracy: 0.6940 - val\_precision: 0.8012 - val\_recall: 0.5930

Epoch 4/50

1562/1562 [=====] - 32s 20ms/step - loss: 0.9584 - accuracy: 0.6708 - precision: 0.7929 - recall: 0.5524 - val\_loss: 0.9401 - val\_accuracy: 0.6887 - val\_precision: 0.7832 - val\_recall: 0.6193

Epoch 5/50

1562/1562 [=====] - 32s 21ms/step - loss: 0.8733 - accuracy: 0.7000 - precision: 0.8000 - recall: 0.5000 - val\_loss: 0.8000 - val\_accuracy: 0.7000 - val\_precision: 0.8000 - val\_recall: 0.5000

026 - precision: 0.8059 - recall: 0.5978 - val\_loss: 1.0669 - val\_accuracy: 0.6685 - val\_precision: 0.7446 - val\_recall: 0.6036  
Epoch 6/50  
1562/1562 [=====] - 32s 20ms/step - loss: 0.8114 - accuracy: 0.7241 - precision: 0.8185 - recall: 0.6331 - val\_loss: 0.8686 - val\_accuracy: 0.7099 - val\_precision: 0.7890 - val\_recall: 0.6427  
Epoch 7/50  
1562/1562 [=====] - 32s 21ms/step - loss: 0.7767 - accuracy: 0.7379 - precision: 0.8264 - recall: 0.6510 - val\_loss: 0.7486 - val\_accuracy: 0.7590 - val\_precision: 0.8239 - val\_recall: 0.6958  
Epoch 8/50  
1562/1562 [=====] - 32s 21ms/step - loss: 0.7301 - accuracy: 0.7546 - precision: 0.8379 - recall: 0.6714 - val\_loss: 0.6864 - val\_accuracy: 0.7717 - val\_precision: 0.8414 - val\_recall: 0.7152  
Epoch 9/50  
1562/1562 [=====] - 32s 21ms/step - loss: 0.6939 - accuracy: 0.7623 - precision: 0.8435 - recall: 0.6887 - val\_loss: 0.6023 - val\_accuracy: 0.7987 - val\_precision: 0.8577 - val\_recall: 0.7477  
Epoch 10/50  
1562/1562 [=====] - 33s 21ms/step - loss: 0.6620 - accuracy: 0.7785 - precision: 0.8550 - recall: 0.7081 - val\_loss: 0.6047 - val\_accuracy: 0.7996 - val\_precision: 0.8687 - val\_recall: 0.7257  
Epoch 11/50  
1562/1562 [=====] - 33s 21ms/step - loss: 0.6272 - accuracy: 0.7892 - precision: 0.8599 - recall: 0.7249 - val\_loss: 0.5915 - val\_accuracy: 0.8092 - val\_precision: 0.8609 - val\_recall: 0.7620  
Epoch 12/50  
1562/1562 [=====] - 33s 21ms/step - loss: 0.6193 - accuracy: 0.7889 - precision: 0.8570 - recall: 0.7253 - val\_loss: 0.5993 - val\_accuracy: 0.8057 - val\_precision: 0.8571 - val\_recall: 0.7566  
Epoch 13/50  
1562/1562 [=====] - 32s 21ms/step - loss: 0.6024 - accuracy: 0.7982 - precision: 0.8646 - recall: 0.7379 - val\_loss: 0.5442 - val\_accuracy: 0.8226 - val\_precision: 0.8766 - val\_recall: 0.7723  
Epoch 14/50  
1562/1562 [=====] - 33s 21ms/step - loss: 0.5830 - accuracy: 0.8039 - precision: 0.8681 - recall: 0.7441 - val\_loss: 0.5739 - val\_accuracy: 0.8097 - val\_precision: 0.8624 - val\_recall: 0.7630  
Epoch 15/50  
1562/1562 [=====] - 33s 21ms/step - loss: 0.5641 - accuracy: 0.8096 - precision: 0.8719 - recall: 0.7556 - val\_loss: 0.6131 - val\_accuracy: 0.8010 - val\_precision: 0.8531 - val\_recall: 0.7588  
Epoch 16/50  
1562/1562 [=====] - 33s 21ms/step - loss: 0.5441 - accuracy: 0.8156 - precision: 0.8736 - recall: 0.7624 - val\_loss: 0.4915 - val\_accuracy: 0.8397 - val\_precision: 0.8824 - val\_recall: 0.8026  
Epoch 17/50  
1562/1562 [=====] - 33s 21ms/step - loss: 0.5356 - accuracy: 0.8186 - precision: 0.8760 - recall: 0.7642 - val\_loss: 0.4851 - val\_accuracy: 0.8414 - val\_precision: 0.8814 - val\_recall: 0.8029  
Epoch 18/50  
1562/1562 [=====] - 33s 21ms/step - loss: 0.5311 - accuracy: 0.8184 - precision: 0.8757 - recall: 0.7671 - val\_loss: 0.4702 - val\_accuracy: 0.8462 - val\_precision: 0.8902 - val\_recall: 0.8039  
Epoch 19/50  
1562/1562 [=====] - 33s 21ms/step - loss: 0.5084 - accuracy: 0.8254 - precision: 0.8791 - recall: 0.7796 - val\_loss: 0.4905 - val\_accuracy: 0.8387 - val\_precision: 0.8781 - val\_recall: 0.8001  
Epoch 20/50  
1562/1562 [=====] - 32s 21ms/step - loss: 0.4997 - accuracy: 0.8270 - precision: 0.8797 - recall: 0.7823 - val\_loss: 0.5039 - val\_accuracy: 0.8325 - val\_precision: 0.8726 - val\_recall: 0.7985  
Epoch 21/50  
1562/1562 [=====] - 33s 21ms/step - loss: 0.4969 - accuracy: 0.8310 - precision: 0.8831 - recall: 0.7856 - val\_loss: 0.5904 - val\_accuracy: 0.8082 - val\_precision: 0.8554 - val\_recall: 0.7761  
Epoch 22/50  
1562/1562 [=====] - 34s 22ms/step - loss: 0.4795 - accuracy: 0.8345 - precision: 0.8825 - recall: 0.7912 - val\_loss: 0.4307 - val\_accuracy: 0.8564 - val\_precision: 0.8985 - val\_recall: 0.8234  
Epoch 23/50  
1562/1562 [=====] - 34s 22ms/step - loss: 0.4680 - accuracy: 0.8

404 - precision: 0.8889 - recall: 0.7969 - val\_loss: 0.4775 - val\_accuracy: 0.8481 - val\_precision: 0.8863 - val\_recall: 0.8117  
Epoch 24/50  
1562/1562 [=====] - 34s 22ms/step - loss: 0.4734 - accuracy: 0.8403 - precision: 0.8890 - recall: 0.7977 - val\_loss: 0.4535 - val\_accuracy: 0.8480 - val\_precision: 0.8858 - val\_recall: 0.8144  
Epoch 25/50  
1562/1562 [=====] - 34s 22ms/step - loss: 0.4548 - accuracy: 0.8441 - precision: 0.8924 - recall: 0.8036 - val\_loss: 0.4601 - val\_accuracy: 0.8521 - val\_precision: 0.8851 - val\_recall: 0.8187  
Epoch 26/50  
1562/1562 [=====] - 31s 20ms/step - loss: 0.4588 - accuracy: 0.8441 - precision: 0.8905 - recall: 0.8031 - val\_loss: 0.4430 - val\_accuracy: 0.8516 - val\_precision: 0.8906 - val\_recall: 0.8156  
Epoch 27/50  
1562/1562 [=====] - 34s 22ms/step - loss: 0.4534 - accuracy: 0.8456 - precision: 0.8916 - recall: 0.8055 - val\_loss: 0.4414 - val\_accuracy: 0.8556 - val\_precision: 0.8947 - val\_recall: 0.8211  
Epoch 28/50  
1562/1562 [=====] - 32s 20ms/step - loss: 0.4505 - accuracy: 0.8497 - precision: 0.8908 - recall: 0.8082 - val\_loss: 0.4372 - val\_accuracy: 0.8554 - val\_precision: 0.8902 - val\_recall: 0.8305  
Epoch 29/50  
1562/1562 [=====] - 35s 22ms/step - loss: 0.4341 - accuracy: 0.8528 - precision: 0.8952 - recall: 0.8167 - val\_loss: 0.4099 - val\_accuracy: 0.8634 - val\_precision: 0.8940 - val\_recall: 0.8363  
Epoch 30/50  
1562/1562 [=====] - 32s 20ms/step - loss: 0.4221 - accuracy: 0.8576 - precision: 0.8995 - recall: 0.8200 - val\_loss: 0.4433 - val\_accuracy: 0.8519 - val\_precision: 0.8911 - val\_recall: 0.8185  
Epoch 31/50  
1562/1562 [=====] - 35s 22ms/step - loss: 0.4292 - accuracy: 0.8531 - precision: 0.8949 - recall: 0.8161 - val\_loss: 0.4755 - val\_accuracy: 0.8442 - val\_precision: 0.8779 - val\_recall: 0.8156  
Epoch 32/50  
1562/1562 [=====] - 32s 20ms/step - loss: 0.4178 - accuracy: 0.8574 - precision: 0.8978 - recall: 0.8209 - val\_loss: 0.4062 - val\_accuracy: 0.8637 - val\_precision: 0.8967 - val\_recall: 0.8355  
Epoch 33/50  
1562/1562 [=====] - 36s 23ms/step - loss: 0.4177 - accuracy: 0.8563 - precision: 0.8987 - recall: 0.8228 - val\_loss: 0.4109 - val\_accuracy: 0.8617 - val\_precision: 0.8983 - val\_recall: 0.8269  
Epoch 34/50  
1562/1562 [=====] - 32s 20ms/step - loss: 0.4112 - accuracy: 0.8588 - precision: 0.8978 - recall: 0.8223 - val\_loss: 0.4021 - val\_accuracy: 0.8683 - val\_precision: 0.8951 - val\_recall: 0.8432  
Epoch 35/50  
1562/1562 [=====] - 36s 23ms/step - loss: 0.4034 - accuracy: 0.8622 - precision: 0.9011 - recall: 0.8274 - val\_loss: 0.3975 - val\_accuracy: 0.8700 - val\_precision: 0.8986 - val\_recall: 0.8450  
Epoch 36/50  
1562/1562 [=====] - 32s 20ms/step - loss: 0.3974 - accuracy: 0.8649 - precision: 0.9007 - recall: 0.8313 - val\_loss: 0.4161 - val\_accuracy: 0.8640 - val\_precision: 0.8939 - val\_recall: 0.8347  
Epoch 37/50  
1562/1562 [=====] - 34s 22ms/step - loss: 0.3964 - accuracy: 0.8654 - precision: 0.9020 - recall: 0.8316 - val\_loss: 0.4171 - val\_accuracy: 0.8598 - val\_precision: 0.8952 - val\_recall: 0.8333  
Epoch 38/50  
1562/1562 [=====] - 34s 22ms/step - loss: 0.3894 - accuracy: 0.8676 - precision: 0.9029 - recall: 0.8365 - val\_loss: 0.3914 - val\_accuracy: 0.8723 - val\_precision: 0.9048 - val\_recall: 0.8468  
Epoch 39/50  
1562/1562 [=====] - 31s 20ms/step - loss: 0.3907 - accuracy: 0.8645 - precision: 0.9020 - recall: 0.8312 - val\_loss: 0.4384 - val\_accuracy: 0.8659 - val\_precision: 0.8925 - val\_recall: 0.8397  
Epoch 40/50  
1562/1562 [=====] - 36s 23ms/step - loss: 0.3889 - accuracy: 0.8656 - precision: 0.9027 - recall: 0.8332 - val\_loss: 0.4334 - val\_accuracy: 0.8619 - val\_precision: 0.8967 - val\_recall: 0.8316  
Epoch 41/50  
1562/1562 [=====] - 32s 20ms/step - loss: 0.3886 - accuracy: 0.8661 - precision: 0.9027 - recall: 0.8332 - val\_loss: 0.4334 - val\_accuracy: 0.8619 - val\_precision: 0.8967 - val\_recall: 0.8316



```

661 - precision: 0.9023 - recall: 0.8336 - val_loss: 0.4122 - val_accuracy: 0.8633 - val_
precision: 0.8934 - val_recall: 0.8362
Epoch 42/50
1562/1562 [=====] - 37s 23ms/step - loss: 0.3806 - accuracy: 0.8
689 - precision: 0.9049 - recall: 0.8371 - val_loss: 0.4034 - val_accuracy: 0.8689 - val_
precision: 0.8991 - val_recall: 0.8443
Epoch 43/50
1562/1562 [=====] - 32s 20ms/step - loss: 0.3801 - accuracy: 0.8
700 - precision: 0.9053 - recall: 0.8372 - val_loss: 0.3901 - val_accuracy: 0.8740 - val_
precision: 0.9030 - val_recall: 0.8504
Epoch 44/50
1562/1562 [=====] - 37s 24ms/step - loss: 0.3795 - accuracy: 0.8
700 - precision: 0.9045 - recall: 0.8391 - val_loss: 0.4225 - val_accuracy: 0.8619 - val_
precision: 0.8942 - val_recall: 0.8421
Epoch 45/50
1562/1562 [=====] - 32s 20ms/step - loss: 0.3709 - accuracy: 0.8
753 - precision: 0.9078 - recall: 0.8456 - val_loss: 0.4008 - val_accuracy: 0.8679 - val_
precision: 0.8966 - val_recall: 0.8468
Epoch 46/50
1562/1562 [=====] - 37s 24ms/step - loss: 0.3662 - accuracy: 0.8
751 - precision: 0.9103 - recall: 0.8459 - val_loss: 0.3810 - val_accuracy: 0.8723 - val_
precision: 0.9038 - val_recall: 0.8472
Epoch 47/50
1562/1562 [=====] - 31s 20ms/step - loss: 0.3695 - accuracy: 0.8
718 - precision: 0.9048 - recall: 0.8428 - val_loss: 0.3763 - val_accuracy: 0.8804 - val_
precision: 0.9034 - val_recall: 0.8606
Epoch 48/50
1562/1562 [=====] - 37s 24ms/step - loss: 0.3750 - accuracy: 0.8
710 - precision: 0.9062 - recall: 0.8416 - val_loss: 0.3982 - val_accuracy: 0.8713 - val_
precision: 0.9012 - val_recall: 0.8471
Epoch 49/50
1562/1562 [=====] - 31s 20ms/step - loss: 0.3674 - accuracy: 0.8
744 - precision: 0.9081 - recall: 0.8451 - val_loss: 0.3881 - val_accuracy: 0.8754 - val_
precision: 0.9010 - val_recall: 0.8553
Epoch 50/50
1562/1562 [=====] - 34s 21ms/step - loss: 0.3590 - accuracy: 0.8
748 - precision: 0.9087 - recall: 0.8460 - val_loss: 0.4086 - val_accuracy: 0.8699 - val_
precision: 0.8963 - val_recall: 0.8476

```

## Model Evaluation

In [14]:

```

plt.figure(figsize=(12, 16))

plt.subplot(4, 2, 1)
plt.plot(r.history['loss'], label='Loss')
plt.plot(r.history['val_loss'], label='val_Loss')
plt.title('Loss Function Evolution')
plt.legend()

plt.subplot(4, 2, 2)
plt.plot(r.history['accuracy'], label='accuracy')
plt.plot(r.history['val_accuracy'], label='val_accuracy')
plt.title('Accuracy Function Evolution')
plt.legend()

plt.subplot(4, 2, 3)
plt.plot(r.history['precision'], label='precision')
plt.plot(r.history['val_precision'], label='val_precision')
plt.title('Precision Function Evolution')
plt.legend()

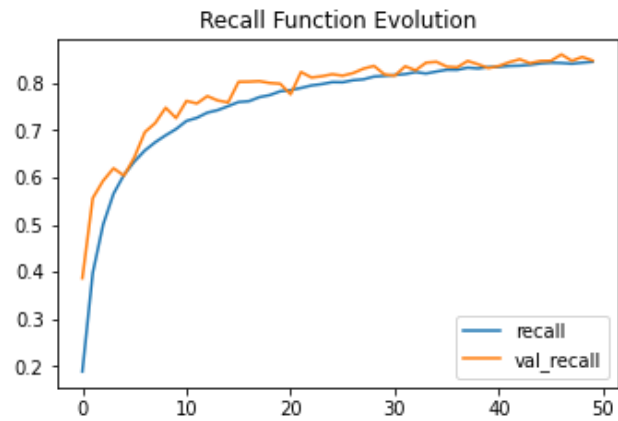


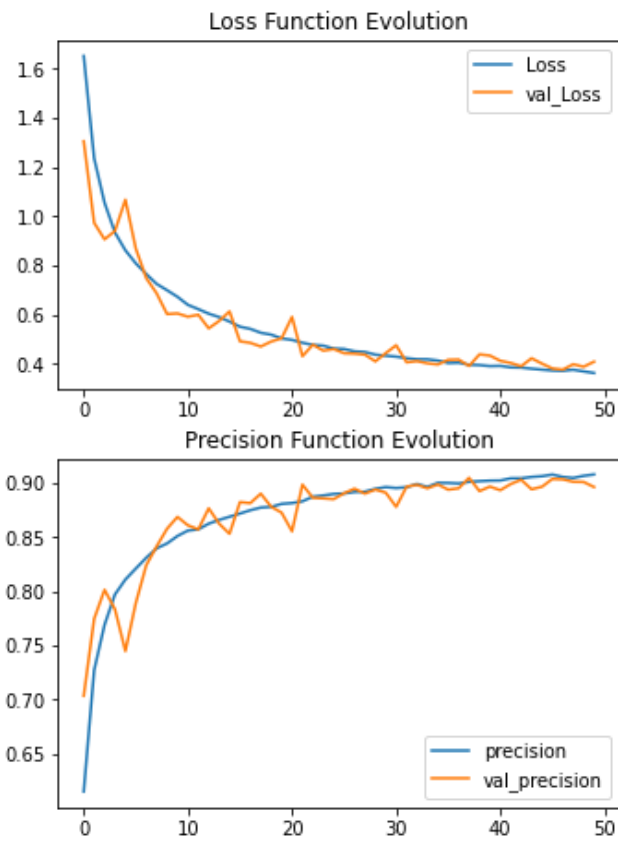
plt.subplot(4, 2, 4)
plt.plot(r.history['recall'], label='recall')
plt.plot(r.history['val_recall'], label='val_recall')
plt.title('Recall Function Evolution')
plt.legend()

```

Out[14]:

<matplotlib.legend.Legend at 0x7f3660724ad0>





In [15]:

```
evaluation = model.evaluate(X_test, y_cat_test)
print(f'Test Accuracy : {evaluation[1] * 100:.2f}%')
```

313/313 [=====] - 1s 5ms/step - loss: 0.4086 - accuracy: 0.8699  
- precision: 0.8963 - recall: 0.8476  
Test Accuracy : 86.99%

In [16]:

```
y_pred = model.predict_classes(X_test)
cm = confusion_matrix(y_test, y_pred)
```


/opt/conda/lib/python3.7/site-packages/tensorflow/python/keras/engine/sequential.py:450:  
UserWarning: `model.predict\_classes()` is deprecated and will be removed after 2021-01-01.  
Please use instead: \* `np.argmax(model.predict(x), axis=-1)`, if your model does multi-class classification (e.g. if it uses a `softmax` last-layer activation). \* `(model.predict(x) > 0.5).astype("int32")`, if your model does binary classification (e.g. if it uses a `sigmoid` last-layer activation).  
warnings.warn("`model.predict\_classes()` is deprecated and ")

In [17]:

```
disp = ConfusionMatrixDisplay(confusion_matrix=cm,
                              display_labels=labels)
```

```
# NOTE: Fill all variables here with default values of the plot_confusion_matrix
fig, ax = plt.subplots(figsize=(10, 10))
disp = disp.plot(xticks_rotation='vertical', ax=ax, cmap='summer')
plt.show()
```





In [18]:

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.89	0.90	0.89	1000
1	0.93	0.95	0.94	1000
2	0.85	0.81	0.83	1000
3	0.77	0.76	0.76	1000
4	0.90	0.81	0.85	1000
5	0.90	0.74	0.81	1000
6	0.76	0.96	0.85	1000
7	0.90	0.92	0.91	1000
8	0.94	0.91	0.93	1000
9	0.88	0.95	0.91	1000
accuracy			0.87	10000
macro avg	0.87	0.87	0.87	10000
weighted avg	0.87	0.87	0.87	10000


MODEL TESTING


In [19]:

```
my_image = X_test[100]
plt.imshow(my_image)
```

Out[19]:

<matplotlib.image.AxesImage at 0x7f3660454610>





In [20]:

```
#DEER
y_test[100]
```

Out[20]:

```
array([4], dtype=uint8)
```

In [21]:

```
# correctly predicted as a Deer
model.predict_classes(my_image.reshape(1, 32, 32, 3))
```

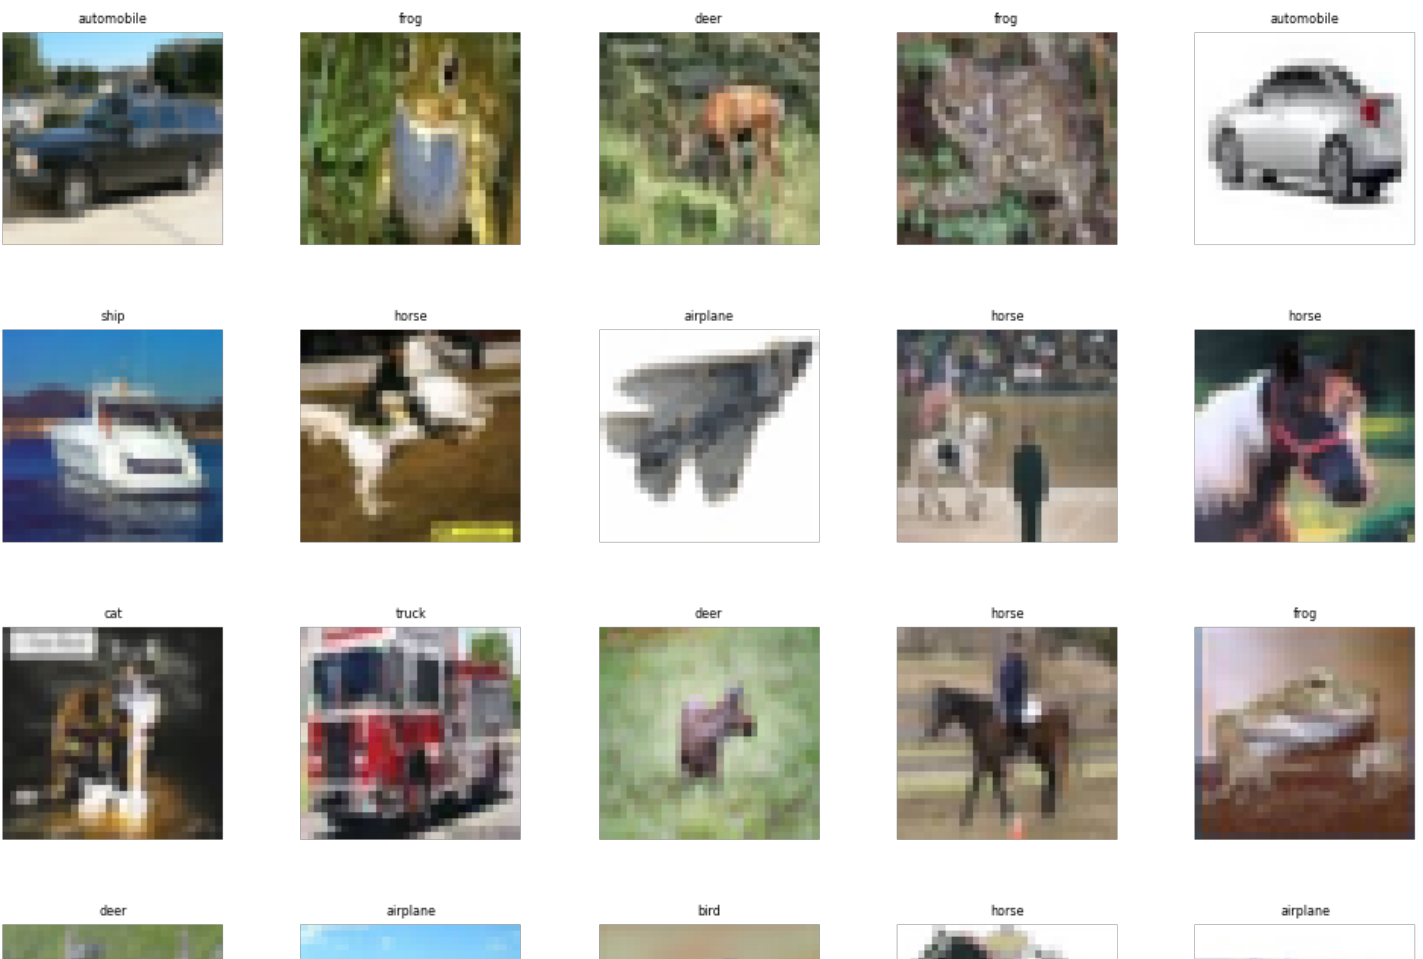
Out[21]:


```
array([7])
```

In [22]:

```
labels = ['airplane', 'automobile', 'bird', 'cat', 'deer',
          'dog', 'frog', 'horse', 'ship', 'truck']
W_grid = 5
L_grid = 5
fig, axes = plt.subplots(L_grid, W_grid, figsize = (17,17))
axes = axes.ravel()
n_test = len(X_test)
for i in np.arange(0, W_grid * L_grid): # create evenly spaces variables
    index = np.random.randint(0, n_test)
    axes[i].imshow(X_test[index,1:])
    label_index = int(y_pred[index])
    axes[i].set_title(labels[label_index], fontsize = 8)
    axes[i].axis('off')

plt.subplots_adjust(hspace=0.4)
```





cat

truck

cat

dog

frog



In [23]:

```
predictions = model.predict(X_test)
```

In [24]:

```
def plot_image(i, predictions_array, true_label, img):
    predictions_array, true_label, img = predictions_array, true_label[i], img[i]
    plt.grid(False)
    plt.xticks([])
    plt.yticks([])

    plt.imshow(img, cmap=plt.cm.binary)

    predicted_label = np.argmax(predictions_array)
    if predicted_label == true_label:
        color = 'blue'
    else:
        color = 'red'





    plt.xlabel(f"{labels[int(predicted_label)]} {100*np.max(predictions_array):2.0f}% ({100*np.max(true_labels[int(true_label)])})",
               color=color)

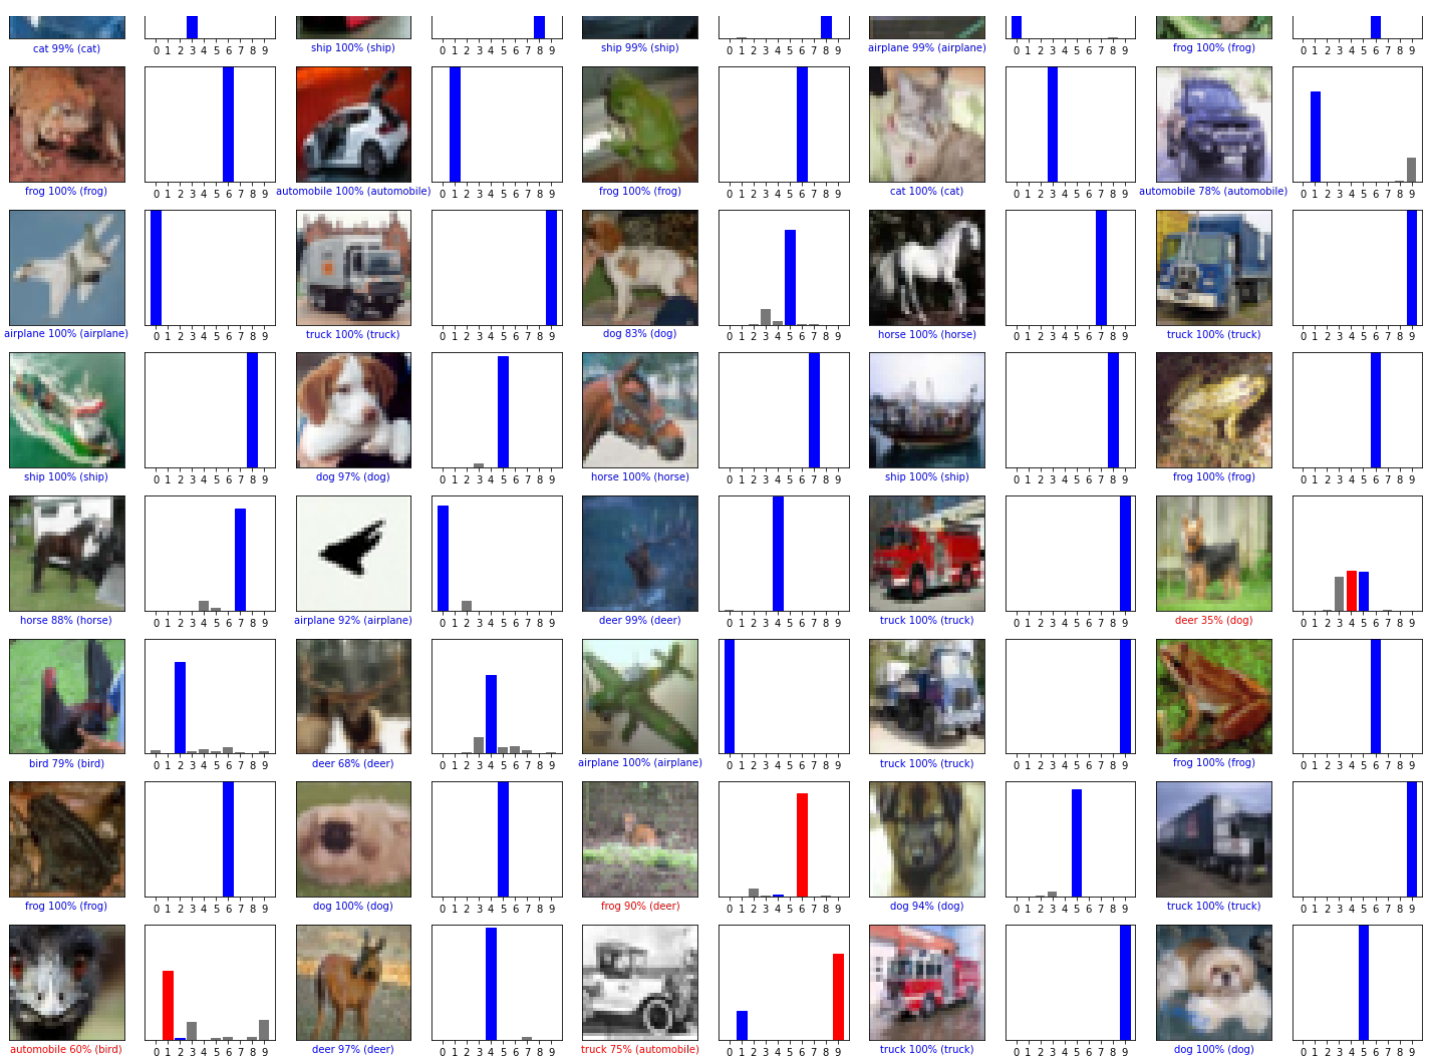
def plot_value_array(i, predictions_array, true_label):
    predictions_array, true_label = predictions_array, int(true_label[i])
    plt.grid(False)
    plt.xticks(range(10))
    plt.yticks([])
    thisplot = plt.bar(range(10), predictions_array, color="#777777")
    plt.ylim([0, 1])
    predicted_label = np.argmax(predictions_array)

    thisplot[predicted_label].set_color('red')
    thisplot[true_label].set_color('blue')
```

In [25]:

```
# Plot the first X test images, their predicted labels, and the true labels.
# Color correct predictions in blue and incorrect predictions in red.
num_rows = 8
num_cols = 5
num_images = num_rows * num_cols
plt.figure(figsize=(2 * 2 * num_cols, 2 * num_rows))
for i in range(num_images):
    plt.subplot(num_rows, 2 * num_cols, 2 * i + 1)
    plot_image(i, predictions[i], y_test, X_test)
    plt.subplot(num_rows, 2*num_cols, 2*i+2)
    plot_value_array(i, predictions[i], y_test)
plt.tight_layout()
plt.show()
```





## Save the models

In [27]:

```
from tensorflow.keras.models import load_model

model.save('Assignment3.h5')
```